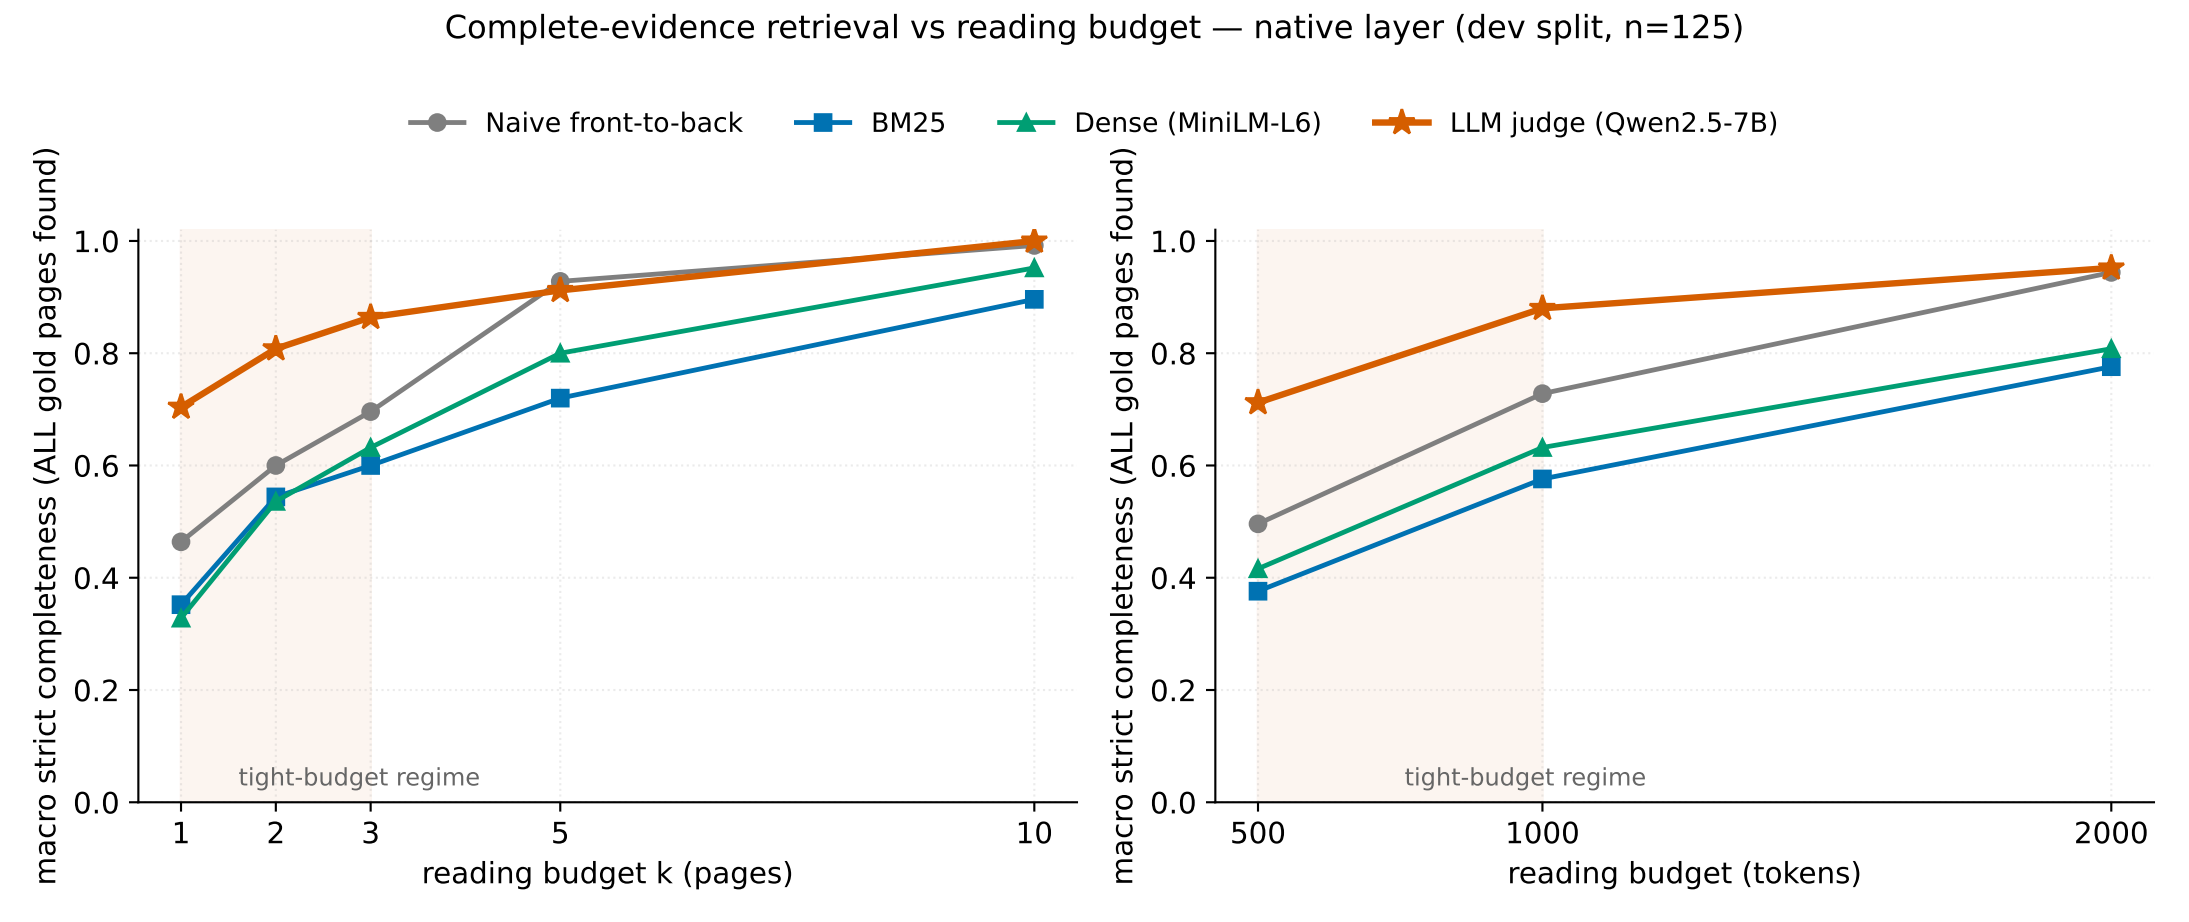
Complete-evidence retrieval vs reading budget — native layer (dev split, n=125)
Naive front-to-back
BM25
Dense (MiniLM-L6)
LLM judge (Qwen2.5-7B)
macro strict completeness (ALL gold pages found)
1.0
0.8
0.6
0.4
0.2
0.0
tight-budget regime
1
2
3
5
10
reading budget k (pages)
macro strict completeness (ALL gold pages found)
tight-budget regime
500
1000
2000
reading budget (tokens)