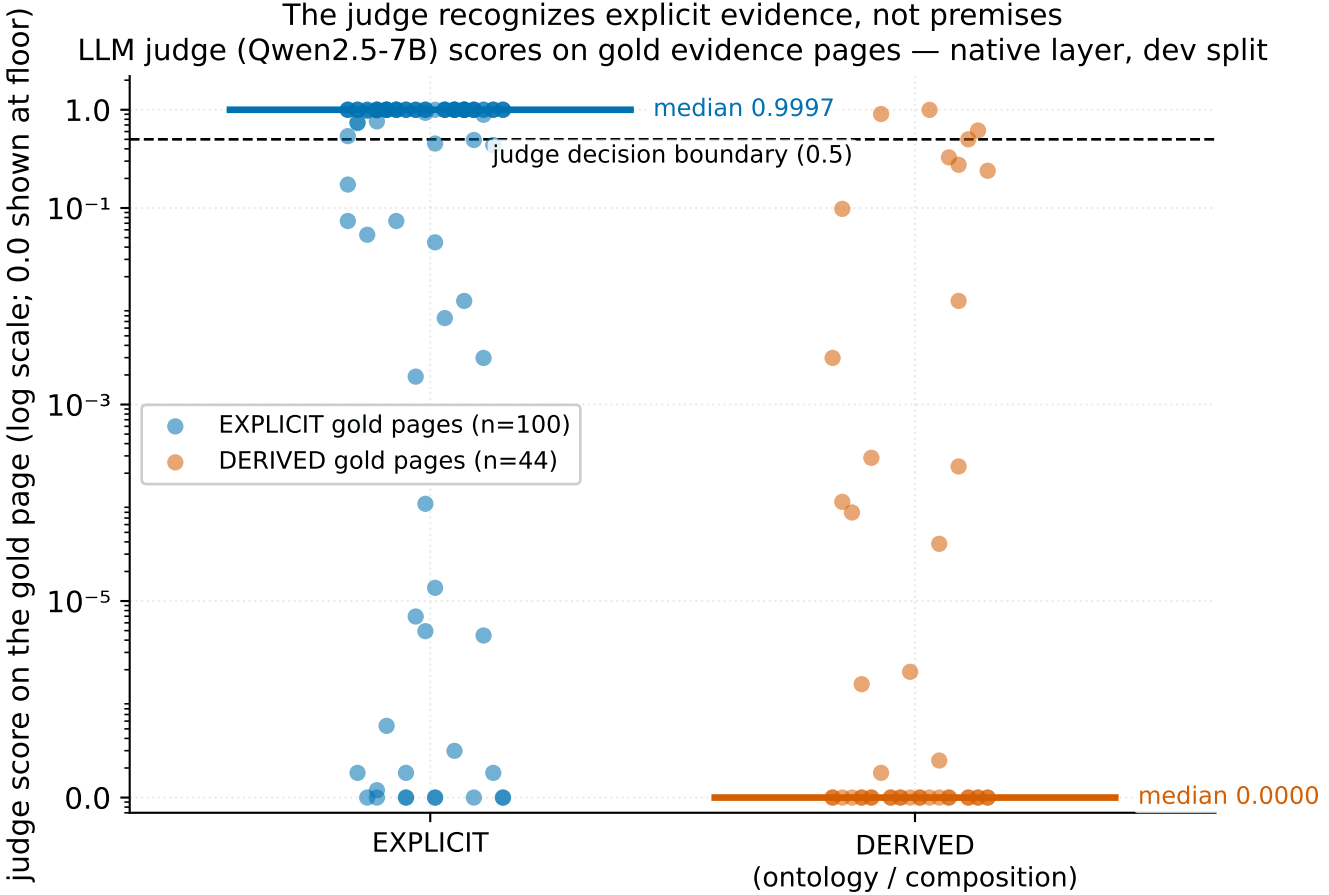
The judge recognizes explicit evidence, not premises
LLM judge (Qwen2.5-7B) scores on gold evidence pages — native layer, dev split
judge score on the gold page (log scale; 0.0 shown at floor)
1.0
10^{-1}
10^{-3}
10^{-5}
0.0
median 0.9997
judge decision boundary (0.5)
EXPLICIT gold pages (n=100)
DERIVED gold pages (n=44)
median 0.0000
EXPLICIT
DERIVED
(ontology / composition)